

Kubernetes on Azure

Scott Coulton
Developer Advocate



@scottcoulton

About me.

Scott Coulton
Developer Advocate

Spent the last 4 years on
container related
development
I love golang
I am also a Docker Captain



@scottcoulton



scotty-c



Agenda

- Kubernetes 101
 - Introduction into Kubernetes
 - Kubernetes components
 - Deploying Kubernetes on Azure
 - Pods, services and deployments
 - Rabc, roles and service accounts
 - Stateful sets
 - Kubernetes networking and service discovery
 - Load balancing and ingress control



Agenda

- Helm
 - Introduction into Helm
 - Understanding charts
 - Deploying Helm on Kubernetes
 - Helm cli
 - Deploying a public chart
 - Writing our own chart
 - Helm and CNAB



Agenda

- Kubernetes advanced topics
 - Virtual node with virtual kubelet
 - Pod security context
 - Introduction to istio
 - Advanced application routing with istio
 - Setting mTLS between application services with istio



@scottcoulton

Course assumptions

Prior knowledge

- A basic understanding of Linux
- Be able to read bash scripts
- Understand what a container is

Equipment needed

- A bash shell (WSL is fine)
- An Azure account with access to create resources service principals
- Azure cli 2.0



Tools we will need

Please install

- kubectl <https://kubernetes.io/docs/tasks/tools/install-kubectl/>
- kubectx <https://github.com/ahmetb/kubectx>
- jq (from your package manager)



Code examples

Code for this course can be downloaded from

<https://github.com/scotty-c/kubernetes-on-azure-workshop>



Introduction into Kubernetes



@scottcoulton

So what is Kubernetes

Kubernetes is a portable, extensible open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available. Google open-sourced the Kubernetes project in 2014.

Why do I need Kubernetes and what can it do?

Kubernetes has a number of features. It can be thought of as:

- a container platform
- a microservices platform
- a portable cloud platform and a lot more

Why do I need Kubernetes and what can it do?

Kubernetes provides a **container-centric** management environment. It orchestrates computing, networking, and storage infrastructure on behalf of user workloads. This provides much of the simplicity of Platform as a Service (PaaS) with the flexibility of Infrastructure as a Service (IaaS), and enables portability across infrastructure providers.

Kubernetes components



@scottcoulton

Kubernetes components

Kubernetes is broken down into two node types.

- Master node
- Worker node



Master node

A master node is responsible for

- Running the control plane
- Scheduling workloads
- Security controls



Worker node

A worker node is responsible for

- Running workloads



Master node

A master nodes components (control plane)

- kube-apiserver
- etcd
- kube-scheduler
- kube-controller-manager
- cloud-controller-manager



Worker node

A worker nodes components

- kubelet
- Kube-proxy



@scottcoulton

Kube-apiserver

The kube-apiserver is responsible for

- The entry point into the cluster
- It exposes the Kubernetes API
- It's a REST service
- Validates and configures data for the api objects

etcd

Consistent and highly-available key value store used as
Kubernetes' backing store for all cluster data

Exciting news Cosmos DB has an etcd api

kube-scheduler

Kube-scheduler is responsible for

- watches newly created pods that have no node assigned, and selects a node for them to run on

Factors taken into account for scheduling decisions include individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference and deadlines

kube-controller-manager

Kube-controller-manager is responsible for

- Node Controller: Responsible for noticing and responding when nodes go down.
- Replication Controller: Responsible for maintaining the correct number of pods for every replication controller object in the system.
- Endpoints Controller: Populates the Endpoints object (that is, joins Services & Pods)
- Service Account & Token Controllers: Create default accounts and API access tokens for new namespaces.

cloud-controller-manager

Cloud-controller-manager is responsible for

- For checking the cloud provider to determine if a node has been deleted in the cloud after it stops responding
- For setting up routes in the underlying cloud infrastructure
- For creating, updating and deleting cloud provider load balancers
- For creating, attaching, and mounting volumes, and interacting with the cloud provider to orchestrate volumes

kubelet

Kubelet is responsible for

- All containers in a pod are running

The kubelet takes a set of PodSpecs that are provided through various mechanisms and ensures that the containers described in those PodSpecs are running and healthy

Kube-proxy

This reflects services as defined in the Kubernetes API on each node and can do simple TCP, UDP, and SCTP stream forwarding or round robin TCP, UDP, and SCTP forwarding across a set of backends

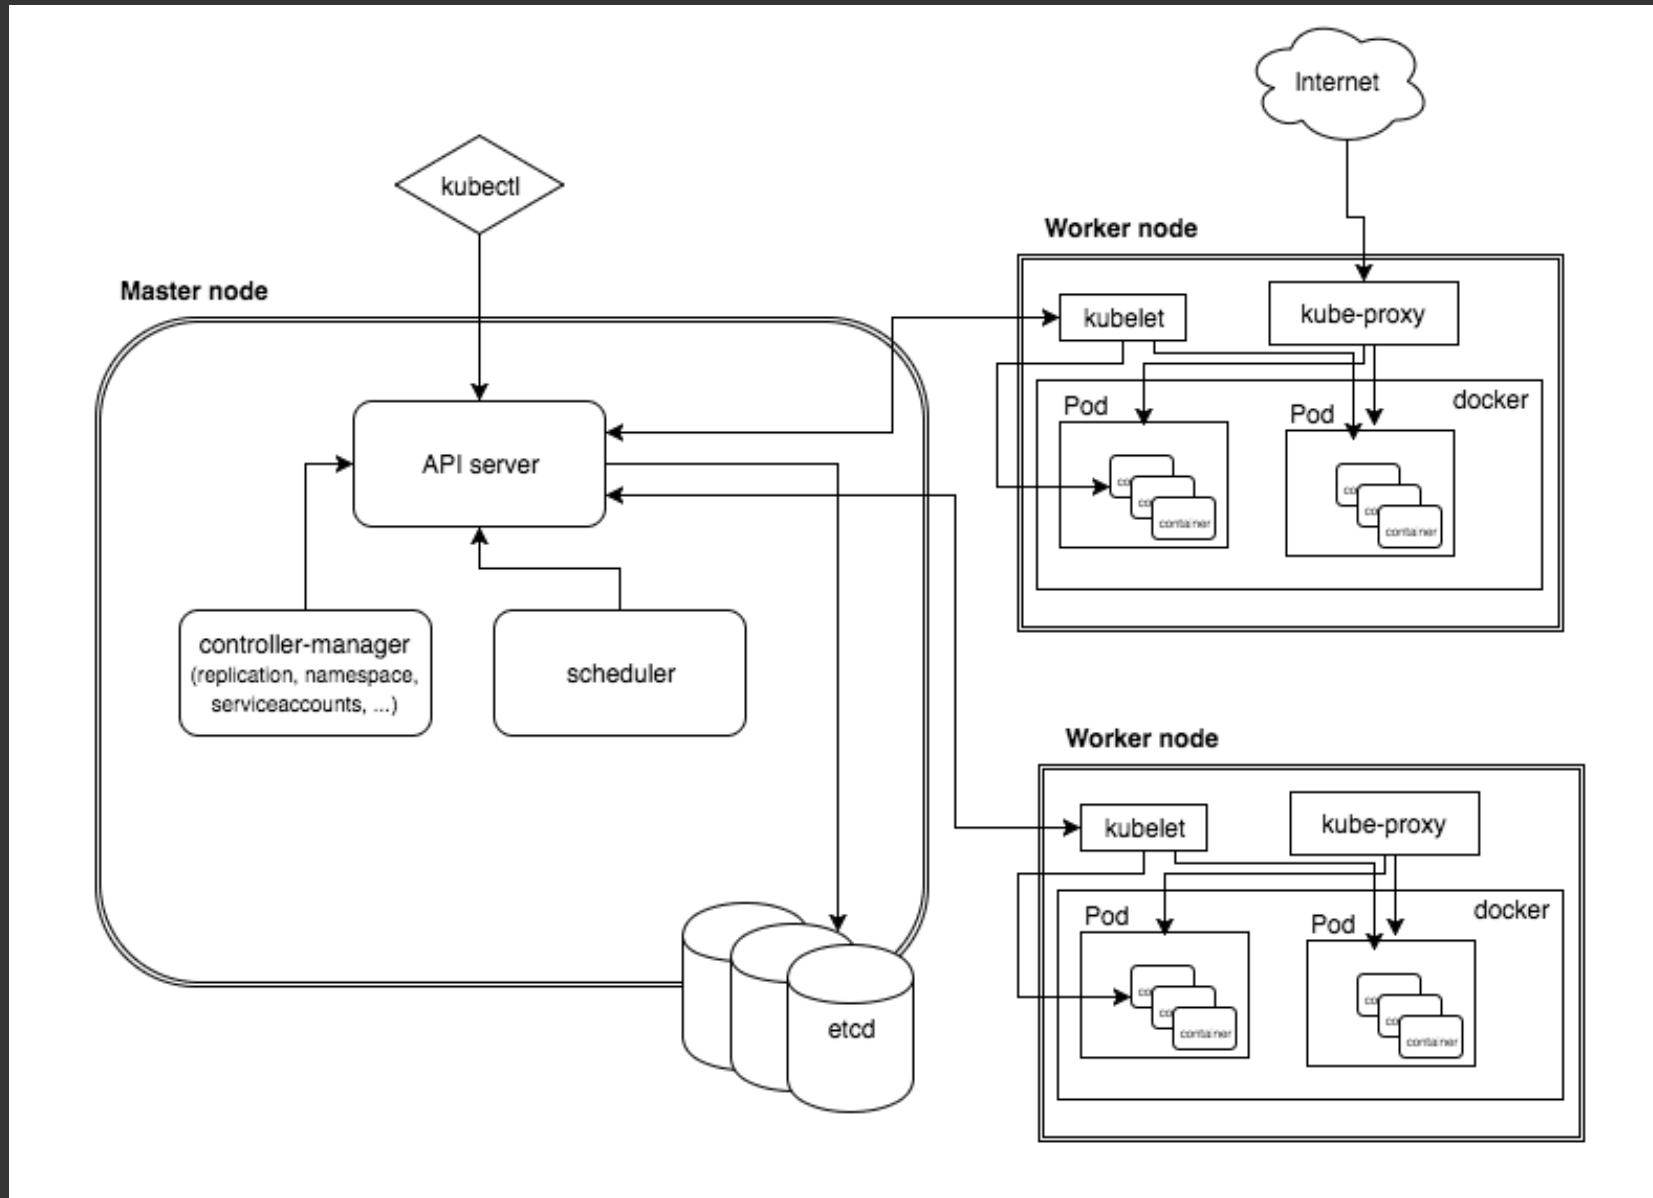
Container runtimes

Kubernetes can use different container runtimes

- Docker
- Moby
- Containerd
- Cri-o

At Azure we use Moby

Kubernetes architecture



Deploying Kubernetes on Azure



@scottcoulton

The offical docs are here

<https://docs.microsoft.com/en-us/azure/aks/kubernetes-walkthrough>

@scottcoulton

Create a resource group

```
az group create --name k8s --location eastus
```

Create your cluster

```
az aks create --resource-group k8s \
  --name k8s \
  --generate-ssh-keys \
  --kubernetes-version 1.12.5 \
  --enable-rbac \
  --node-vm-size Standard_DS2_v2
```

If you don't have kubectl already

```
az aks install-cli
```

@scottcoulton

Get cluster credentials

```
az aks get-credentials --resource-group k8s -  
-name k8s
```

Test out your cluster

```
kubectl get nodes
```

```
kubectl get pods --all-namespaces
```

Pods, services and deployments



@scottcoulton

“A pod is not equal to a
container”

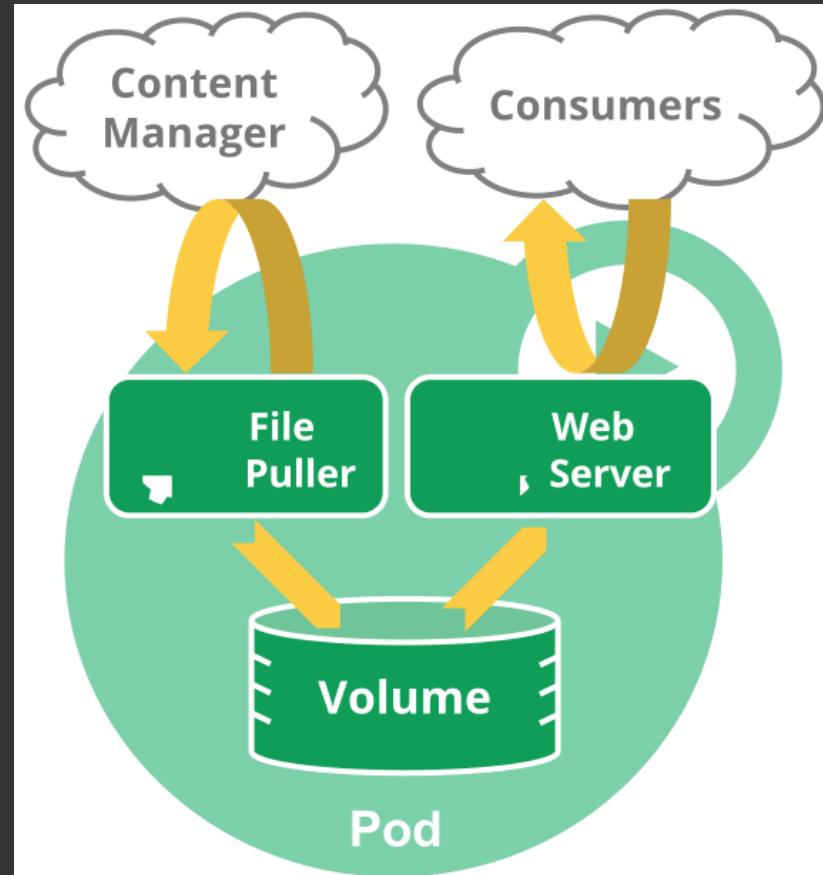
Pods

Pods are a single or group of containers that share

- localhost
- storage
- ip address
- port range

The shared context of a pod is a set of Linux namespaces, cgroups, and potentially other facets of isolation

Pods



@scottcoulton

Defining a pod

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
    - name: myapp-container
      image: busybox
        command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

Services

A service in Kubernetes exposes a set of pods

- Creates a vip
- Sets up basic routing to the pods
- Talks to the cloud-manager to assign a public ip or load balancer

Defining a service

```
kind: Service
apiVersion: v1
metadata:
  name: my-service
spec:
  selector:
    app: MyApp
  ports:
  - protocol: TCP
    port: 80
    targetPort: 9376
```

Replicaset

A replicaset defines that state of a running application

- The amount of pods that should are running
- It self-heals the pods to their disered state

Deployments

A deployment defines the lifecycle of an application

- Is made up of pods
- It controls replicsets
- Includes the functionality to update the desired state
- Rolling updates are included

Defining a deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.9
          ports:
            - containerPort: 80
```

Let's deploy our own deployment



@scottcoulton

Our deployment

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 3
  template:
    metadata:
      labels:
        app: webapp
  spec:
    containers:
      - name: webapp
        image: scottyc/webapp:latest
    ports:
      - containerPort: 3000
        hostPort: 3000
EOF
```

Check our deployment

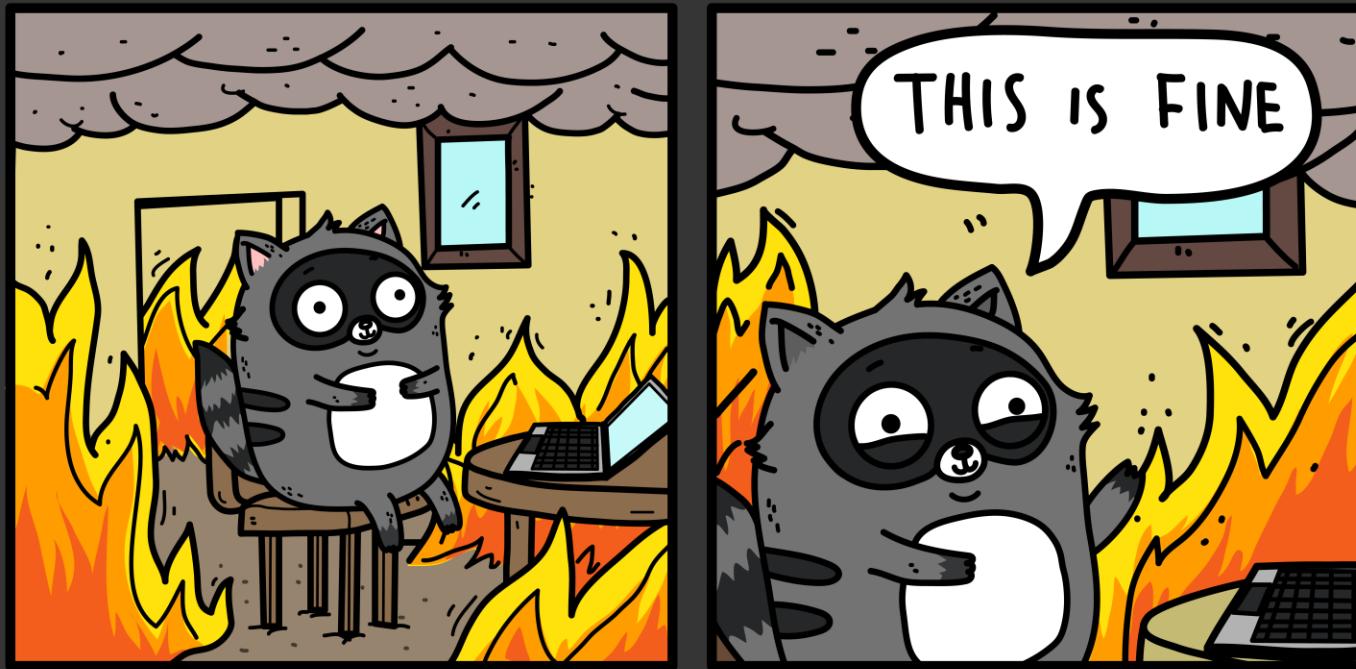
`kubectl get deployments`

`kubectl get pods`

`kubectl get service`

@scottcoulton

What issue did you find ?



@scottcoulton

Answer: We have not exposed a service



@scottcoulton

Exposing our deployment

```
kubectl expose deployment webapp-deployment --  
type=LoadBalancer
```

```
kubectl get service
```

To access our app <http://<your-pub-ip>:3000>



@scottcoulton

Rbac, roles and service accounts



@scottcoulton

rbac

Role based access control (rbac)

- Separation of applications
- Access control for users
- Access control for applications (service accounts)

namespaces

Namespaces are the logical separation in kubernetes

Things that are namespaced

- dns <service-name>.<namespace-name>.svc.cluster.local
- Deployments, services and pods
- Access control for applications (service accounts)
- Resource quotas
- Secrets

namespaces

Things that are NOT namespaced

- Nodes
- Networking
- Storage

Service accounts vs user accounts

The differences are

- User accounts are for humans. Service accounts are for processes, which run in pods.
- User accounts are intended to be global. Names must be unique across all namespaces of a cluster, future user resource will not be namespaced. Service accounts are namespaced.

Create a namespace

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Namespace
metadata:
  name: webapp-namespace
EOF
```

Create a service account

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: ServiceAccount
metadata:
  name: webapp-service-account
  namespace: webapp-namespace
EOF
```

Create a role

```
cat <<EOF | kubectl apply -f -
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: webapp-role
  namespace: webapp-namespace
rules:
- apiGroups: []
  resources: ["pods", "pods/log"]
  verbs: ["get", "list", "watch"]
EOF
```

Create a role binding

```
cat <<EOF | kubectl apply -f -
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: webapp-role-binding
  namespace: webapp-namespace
subjects:
- kind: ServiceAccount
  name: webapp-service-account
  namespace: webapp-namespace
roleRef:
  kind: Role
  name: webapp-role
  apiGroup: rbac.authorization.k8s.io
EOF
```

Deploying an application to our namespace

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp-deployment
  namespace: webapp-namespace
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
          ports:
            - containerPort: 3000
              hostPort: 3000
EOF
```

@scottcoulton

Now let's set up kubectl to use the service account



@scottcoulton

We need the secret for the service account

```
SECRET_NAME=$(kubectl get sa webapp-service-account --namespace webapp-namespace -o json | jq -r '.secrets[].name')
```

We will get the ca

```
kubectl get secret --namespace webapp-namespace  
"${SECRET_NAME}" -o json | jq -r  
'.data["ca.crt"]' | base64 --decode > ca.crt
```

We will get the user token

```
kubectl get secret --namespace webapp-namespace  
"${SECRET_NAME}" -o json | jq -r  
'.data["ca.crt"]' | base64 --decode > ca.crt
```

@scottcoulton

Then we will create our kubeconfig file

```
context=$(kubectl config current-context)

CLUSTER_NAME=$(kubectl config get-contexts "$context" | awk '{print $3}' | tail -n 1)

ENDPOINT=$(kubectl config view -o jsonpath=".clusters[?(@.name == \"${CLUSTER_NAME}\")].cluster.server")

kubectl config set-cluster "${CLUSTER_NAME}" --kubeconfig=admin.conf --server="${ENDPOINT}" --certificate-authority=ca.crt --embed-certs=true

kubectl config set-credentials "webapp-service-account-webapp-namespace-${CLUSTER_NAME}" --kubeconfig=admin.conf --token="${USER_TOKEN}"

kubectl config set-context "webapp-service-account-webapp-namespace-${CLUSTER_NAME}" --kubeconfig=admin.conf --cluster="${CLUSTER_NAME}" --user="webapp-service-account-webapp-namespace-${CLUSTER_NAME}" --namespace webapp-namespace

kubectl config use-context "webapp-service-account-webapp-namespace-${CLUSTER_NAME}" --kubeconfig="${KUBECFG_FILE_NAME}"
```

Export the file to use in the current terminal

```
export KUBECONFIG=admin.conf
```

@scottcoulton

Use your kubectl commands to see what you have access too



@scottcoulton

Statefull sets



@scottcoulton

Stateful sets

Most applications will need some form of state. This is where stateful sets come in handy. This will allow us to persist our data

Before we can access the volumes we need to set up

- storage class
- pvc (persistent volume claim)

Storage classes

Storage classes are the mechanism to provision storage on various backends.

Some of the common backends are

- Local disk, NFS, iSCSI
- Cloud disks ([Azure disk](#), AWS EBS)
- Advanced replicated storage (Rook, Portworx)

Storage classes Azure disk

In this course we will use the dynamic storage class that ships with AKS



@scottcoulton

Persistent volume claims

Persistent volume claims are the units of storage that can be attached to pods.

PVC are configured into two classes

- static
- dynamic

Creating a static claim

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: aks-volume-claim
spec:
  accessModes:
  - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
EOF
```

@scottcoulton

Using the claim

```
cat <<EOF | kubectl apply -f -
kind: Pod
apiVersion: v1
metadata:
  name: nginx-pvc
spec:
  volumes:
    - name: nginx-storage
      persistentVolumeClaim:
        claimName: aks-volume-claim
  containers:
    - name: task-pv-container
      image: nginx
      ports:
        - containerPort: 80
          name: "http-server"
  volumeMounts:
    - mountPath: "/usr/share/nginx/html"
      name: nginx-storage
EOF
```

@scottcoulton

Ingress controller



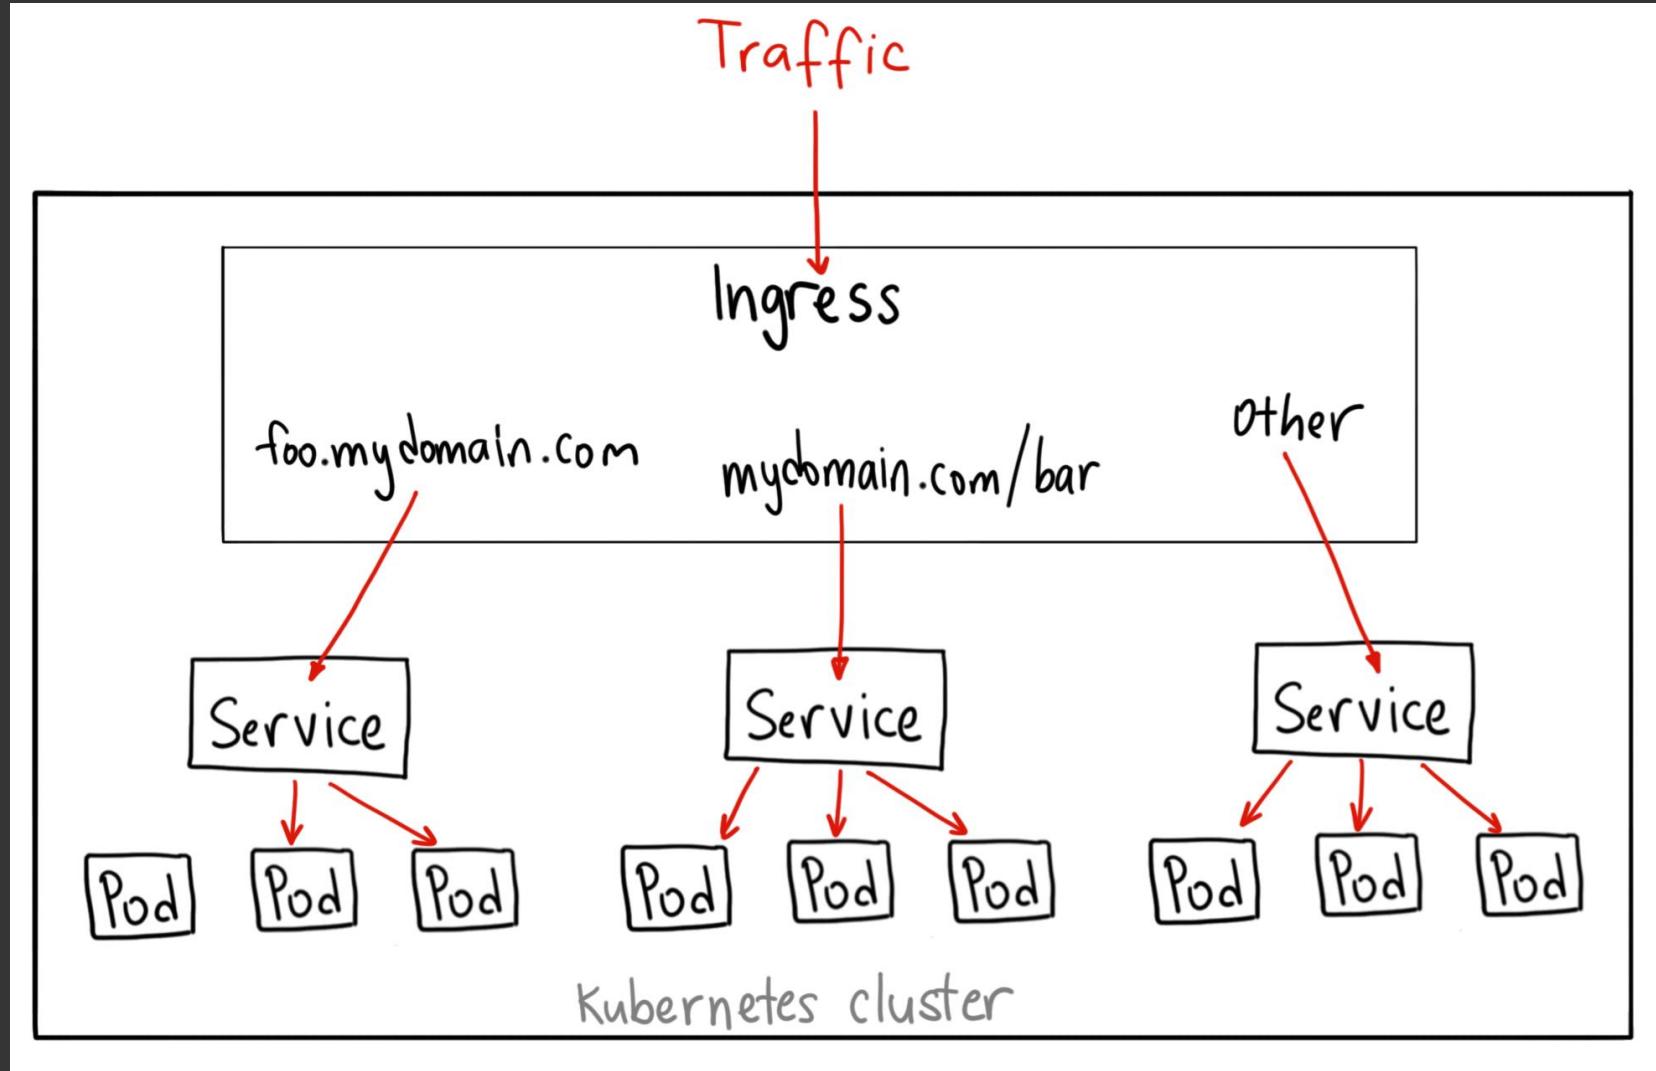
@scottcoulton

Ingress controllers

We have services what do ingress controllers give us.

- Layer 7 routing
- ssl termination
- Single IP address for multiple services

Ingress controllers



Azure HTTP routing addon

Disclaimer !!! This addon is not for production use.

<https://docs.microsoft.com/azure/aks/http-application-routing>

Azure HTTP routing addon

The addon gives us two things straight out of the box.

- A single ingress controller

The Ingress controller is exposed to the internet by using a Kubernetes service of type LoadBalancer. The Ingress controller watches and implements [Kubernetes Ingress resources](#), which creates routes to application endpoints.

- External-DNS controller

Watches for Kubernetes Ingress resources and creates DNS A records in the cluster-specific DNS zone.

Enable the addon

```
az aks enable-addons --resource-group k8s --  
name k8s --addons http_application_routing
```

Deploy our application

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
      ports:
        - containerPort: 3000
          hostPort: 3000
EOF
```

Deploy a service

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Service
metadata:
  name: webapp
spec:
  ports:
  - port: 80
    protocol: TCP
    targetPort: 3000
  selector:
    app: webapp
  type: ClusterIP
EOF
```

Add an ingress rule

```
DNS=$(az aks show --resource-group k8s --name k8s --query  
addonProfiles.httpApplicationRouting.config.HTTPApplicationRoutingZoneName -o tsv)
```

```
cat <<EOF | kubectl apply -f -  
apiVersion: extensions/v1beta1  
kind: Ingress  
metadata:  
  name: webapp  
  annotations:  
    kubernetes.io/ingress.class: addon-http-application-routing  
spec:  
  rules:  
  - host: webapp.$DNS  
    http:  
      paths:  
      - backend:  
          serviceName: webapp  
          servicePort: 80  
        path: /  
EOF
```

Introduction into Helm



@scottcoulton

Introduction into Helm

Helm is a framework to deploy complex applications in Kubernetes.

Helm allows you to tie together

- deployments
- storage
- services

One of the main focuses of Helm is to make your applications shareable and portable

Introduction into Helm

There are 3 main components to Helm.

- Tiller (server side)
- Helm (client side)
- Charts (application definitions)

Introduction into Helm

Tiller is responsible for

- Listening for incoming requests from the Helm client
- Combining a chart and configuration to build a release
- Installing charts into Kubernetes, and then tracking the subsequent release
- Upgrading and uninstalling charts by interacting with Kubernetes

Introduction into Helm

Helm is responsible for

- Local chart development
- Managing repositories
- Interacting with the Tiller server
 - Sending the charts to be installed
 - Asking for information about release
 - Requesting upgrading or uninstalling of existing releases

Introduction into Helm

Helm uses a packaging format called *charts*. A chart is a collection of files that describe a related set of Kubernetes resources.

- Can be a single resource or many
- They use semantic versioning
- Charts can be shared through a chart repository

Introduction into Helm

The official public chart repository lives here.

<https://github.com/helm/charts>

There are two main folders stable and incubation

Understanding charts



@scottcoulton

Understanding charts

For charts to be reusable there needs to be a definition for developers to follow.

We will look at

- Folder structure
- Different files and what they are used for
- Versioning

Folder structure

```
wordpress/
  Chart.yaml          # A YAML file containing information about the chart
  LICENSE            # OPTIONAL: A plain text file containing the license for the chart
  README.md          # OPTIONAL: A human-readable README file
  requirements.yaml # OPTIONAL: A YAML file listing dependencies for the chart
  values.yaml        # The default configuration values for this chart
  charts/             # A directory containing any charts upon which this chart depends.
  templates/          # A directory of templates that, when combined with values,
                     # will generate valid Kubernetes manifest files.
  templates/NOTES.txt # OPTIONAL: A plain text file containing short usage notes
```

Chart.yaml

```
apiVersion: The chart API version, always "v1" (required)
name: The name of the chart (required)
version: A SemVer 2 version (required)
kubeversion: A SemVer range of compatible Kubernetes versions (optional)
description: A single-sentence description of this project (optional)
keywords:
  - A list of keywords about this project (optional)
home: The URL of this project's home page (optional)
sources:
  - A list of URLs to source code for this project (optional)
maintainers: # (optional)
  - name: The maintainer's name (required for each maintainer)
    email: The maintainer's email (optional for each maintainer)
    url: A URL for the maintainer (optional for each maintainer)
engine: gotpl # The name of the template engine (optional, defaults to gotpl)
icon: A URL to an SVG or PNG image to be used as an icon (optional).
appVersion: The version of the app that this contains (optional). This needn't be SemVer.
deprecated: Whether this chart is deprecated (optional, boolean)
tillerVersion: The version of Tiller that this chart requires. This should be expressed as a
SemVer range: ">2.0.0" (optional)
```

Values.yaml

This file allows you to define variables that are injected into the go template at runtime.

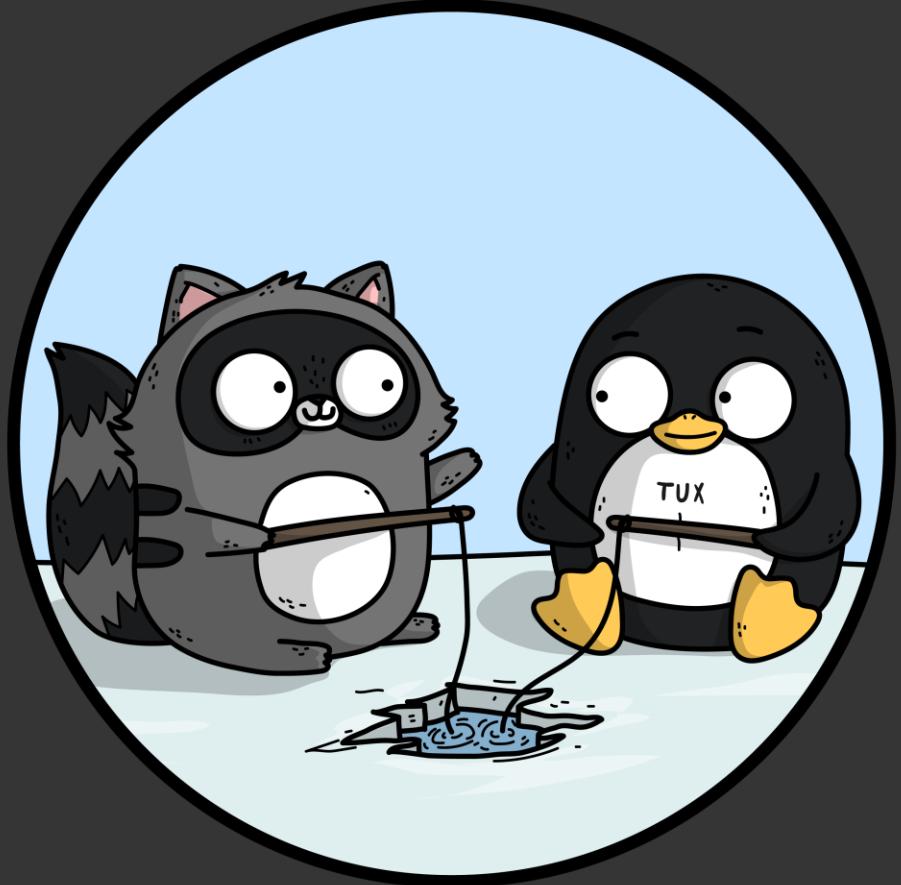
Exmaples would be

- Database connection strings
- Secrets
- Url's

Requirements.yaml

This file allows you to define any other charts that your chart depends on. This works in the same fashion as importing a library in other coding languages.

Installing Helm on Kubernetes



@scottcoulton

The offical docs are here

<https://docs.microsoft.com/azure/aks/kubernetes-helm>

@scottcoulton

Installing Helm

```
#!/bin/bash

if [[ "$OSTYPE" == "linux-gnu" ]]; then
    OS="linux"
    ARCH="linux-amd64"
elif [[ "$OSTYPE" == "darwin"* ]]; then
    OS="osx"
    ARCH="darwin-amd64"
fi
HELM_VERSION=2.11.0

curl -sL "https://storage.googleapis.com/kubernetes-helm/helm-v\$HELM\_VERSION-\$ARCH.tar.gz" | tar xz
chmod +x $ARCH/helm
sudo mv linux-amd64/helm /usr/local/bin/
```

Installing Tiller

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: ServiceAccount
metadata:
  name: tiller
  namespace: kube-system
---
apiVersion: rbac.authorization.k8s.io/v1beta1
kind: ClusterRoleBinding
metadata:
  name: tiller
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: tiller
  namespace: kube-system
EOF
```

@scottcoulton

Installing Tiller

```
helm init --service-account tiller
```

@scottcoulton

Helm cli



@scottcoulton

Helm cli

The Helm cli is the user entry point into Helm

It allows the user to

- Install Tiller into Kubernetes
- Deploy charts
- Preform upgrades of charts
- Add additional chart repos

Helm cli

The Helm cli is the user entry point into Helm

It allows the user to

- Install Tiller into Kubernetes
- Deploy charts
- Preform upgrades of charts
- Add additional chart repos

Handy commands

helm install # install a chart archive

helm delete # given a release name, delete the release from Kubernetes

helm list # list releases

helm fetch # download a chart from a repository

Handy commands

`helm upgrade` # upgrade a release

`helm rollback` # roll back a release to a previous revision

`helm history` # fetch release history

`helm verify` # verify that a chart at the given path has been signed and is valid

`helm create` # creates a new chart

@scottcoulton

Deploying a public chart



@scottcoulton

We are going to deply consul

<https://github.com/helm/charts/tree/master/stable/consul>

@scottcoulton

Why consul ?

Consul is interesting because it's a replicated distributed database.

To work it needs

- Storage
- Networking
- Replication controller

Deploying consul

```
helm install --name consul --set  
StorageClass=default stable/consul
```

To watch your cluster come up

```
kubectl get pods --namespace=default -w
```

Something interesting happened

To find out use the following kubectl command

```
kubectl get pvc
```

Writing our own chart



@scottcoulton

Writing our first chart

Helm has built in tooling to make writing charts easier.

'Helm create' creates

- Our folder structure
- Our templates
- Default values

Writing our first chart

```
helm create mychart
```

Writing our first chart

```
helm install --dry-run --debug ./mychart
```

Writing our first chart

```
helm install --name example ./mychart --set  
service.type=LoadBalancer
```

Writing our first chart

To get our load balancer ip (this might take a minute or so)

```
kubectl get svc --namespace default example-mychart -o jsonpath='{.status.loadBalancer.ingress[0].ip}'
```

Then check that ip address in the browser

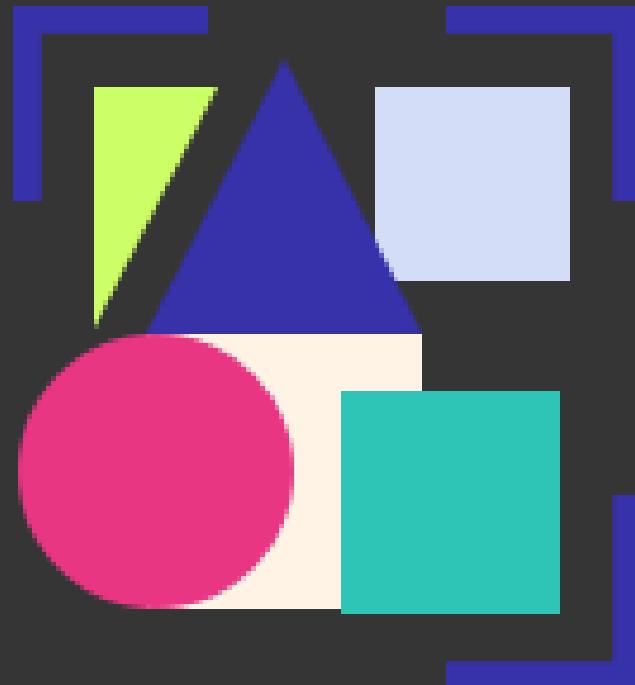
Helm and CNAB



What is CNAB ?

CNAB (Cloud native application bundle) is a specification created by Docker and Microsoft.

<https://cnab.io/>



@scottcoulton

Why do we care about a spec ?



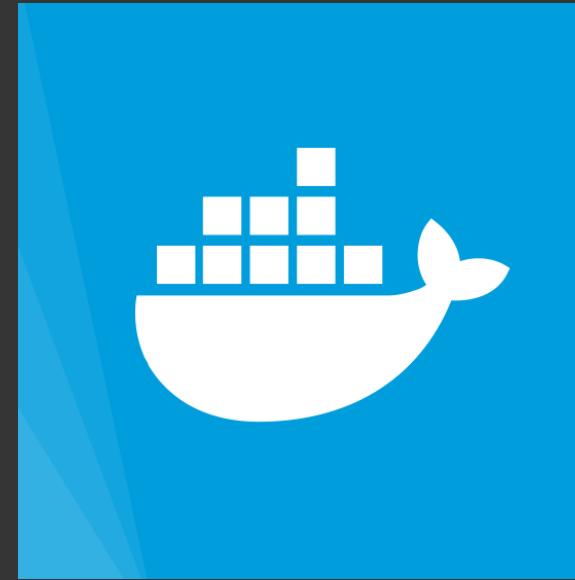
@scottcoulton

What does the spec define ?

- The metadata of the bundle
- The invocation image format
- The bundle runtime
- Claims system
- Security model ie signing

For more reading on the spec head [here](#)

Projects that have implemented the spec



@scottcoulton

Let's look at how duffle interacts with Helm



@scottcoulton

Let's look at how duffle interacts with Helm

<https://github.com/deislabs/example-bundles/tree/master/hellohelm>

@scottcoulton

Virtual node with Virtual Kubelet



What is virtual kubelet ?

Virtual Kubelet is an opensource implementation of kubelet.
This allows kubernetes to schedule pods on platforms that have no nodes

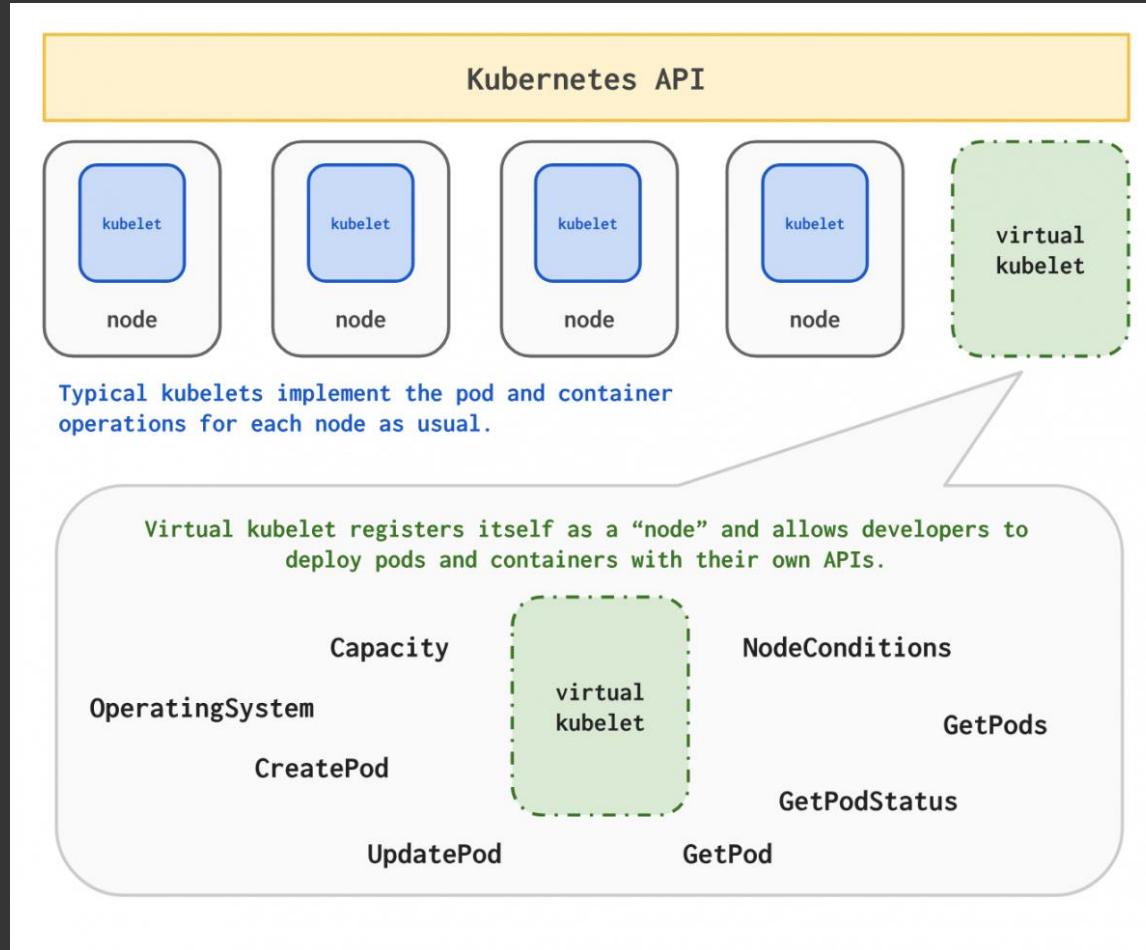
- Azure container service
- AWS Fargate
- Hyper.sh

What does virtual kubelet give us

Virtual Kubelet gives us

- Burstable workloads
- Serverless intergration
- The ability to run Windows and Linux containers
- Cost savings

Virtual kubelet



@scottcoulton

Deploying Virtual Kubelet

```
az aks install-connector --resource-group k8s -  
-name k8s --os-type both
```

Test your Virtual Kubelet

```
kubectl get nodes
```

@scottcoulton

Deploying a pod to virtual kubelet

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: vk-webapp
spec:
  containers:
  - image: scottyc/webapp
    imagePullPolicy: Always
    name: vk-webapp
    resources:
      requests:
        memory: 1G
        cpu: 1
    ports:
    - containerPort: 3000
      name: http
      protocol: TCP
  dnsPolicy: ClusterFirst
  nodeSelector:
    kubernetes.io/role: agent
    beta.kubernetes.io/os: linux
    type: virtual-kubelet
  tolerations:
  - key: virtual-kubelet.io/provider
    operator: Exists
  - key: azure.com/aci
    effect: NoSchedule
EOF
```

Deploying an application on to virtual kubelet

nodeSelector:

 kubernetes.io/role: agent

 beta.kubernetes.io/os: linux

 type: virtual-kubelet

tolerations:

- key: virtual-kubelet.io/provider

- operator: Exists

- key: azure.com/aci

- effect: NoSchedule

Test your Virtual Kubelet

```
kubectl get pods -o wide
```

Then use the public ip to access our application in the browser
http://<public_ip>:3000

Pod security context



@scottcoulton

Pod security context

Just because we are using Kubernetes means we are secure by default.

There are a lot of good security features in Kubernetes that are not turned on.

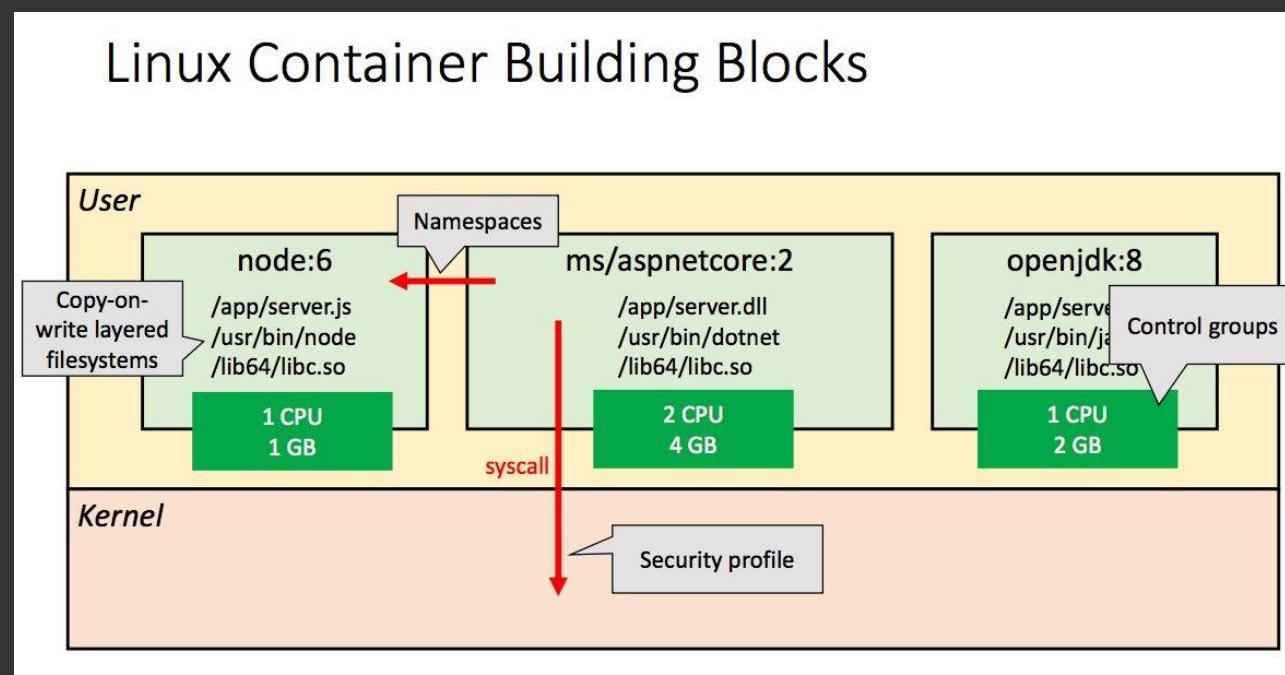
Pod security context

Pod security is an abstraction from the Linux security subsystem.

- Apparmor
- Selinux
- Secomp

Pod security context

A container is a process that is isolated via kernel namespaces and cgroups



@scottcoulton

Pod security context

In Azure our pods are talking to the kernel via Moby

Today we are going to look at three of the important default policies.

In a production environment I would personally use seccomp

https://www.kernel.org/doc/html/v4.16/userspace-api/seccomp_filter.html

With *libseccomp*

<https://github.com/seccomp/libseccomp>

Pod security context

The three default policies are

- runAsUser
- readOnlyRootFilesystem
- allowPrivilegeEscalation

Pod security context

We are going to run a series of deployments with our webapp changing the security context each time.

First, we must see why security context are so important.

Deploy our webapp

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
      ports:
        - containerPort: 3000
          hostPort: 3000
EOF
```

Check your deployment

`kubectl get svc` (to get our public ip)

Check your browser with <http://<public-ip>:3000>

The hack

```
kubectl get pods | grep webapp (to get your pod name)
```

We will then get a shell inside our container
kubectl exec -it <pod_name> sh

Change the index.html file

```
cd static && vim index.html
```

Change the url on line 16

```
to https://media.giphy.com/media/DBfYJqH5AokgM/  
giphy.gif
```

The hack

See what user is the pod running as

whoami from inside the pods terminal

Check your deployment

`kubectl get svc` (to get our public ip)

Check your browser again with <http://<public-ip>:3000>

Make sure you refresh your browser

Defining pod security context

Pod security policy is set in your deployment yaml under:

```
spec:  
  containers:  
    securityContext:
```

Adding runAsUser policy

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
      ports:
        - containerPort: 3000
          hostPort: 3000
      securityContext:
        runAsUser: 1000
EOF
```

@scottcoulton

Check the user running in the container

```
kubectl get pods | grep webapp (to get your pod name)
```

We will then get a shell inside our container

```
kubectl exec -it <pod_name> sh
```

Check the user with

```
whoami
```

Now let's test if we can change the file ? Or change to root

Clean up

```
kubectl delete deployments.apps webapp-deployment
```

Adding readOnlyRootFilesystem policy

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
          ports:
            - containerPort: 3000
              hostPort: 3000
      securityContext:
        readOnlyRootFilesystem: true
EOF
```

@scottcoulton

Check the readonly file system

```
kubectl get pods | grep webapp (to get your pod name)
```

We will then get a shell inside our container

```
kubectl exec -it <pod_name> sh
```

Now let's test if we can change the file ?

Clean up

```
kubectl delete deployments.apps webapp-deployment
```

Adding allowPrivilegeEscalation policy

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
      ports:
        - containerPort: 3000
          hostPort: 3000
      securityContext:
        allowPrivilegeEscalation: false
EOF
```

@scottcoulton

Check the readonly file system

```
kubectl get pods | grep webapp (to get your pod name)
```

We will then get a shell inside our container

```
kubectl exec -it <pod_name> sh
```

Now let's test if we can change the file ? Or change to root

Clean up

```
kubectl delete deployments.apps webapp-deployment
```

Adding all three policies

```
cat <<EOF | kubectl apply -f -
apiVersion: apps/v1 # for versions before 1.9.0 use apps/v1beta2
kind: Deployment
metadata:
  name: webapp-deployment
spec:
  selector:
    matchLabels:
      app: webapp
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:latest
          ports:
            - containerPort: 3000
              hostPort: 3000
      securityContext:
        runAsUser: 1000
        readOnlyRootFilesystem: true
        allowPrivilegeEscalation: false
EOF
```

@scottcoulton

Check with all policies applied

Now let's run all our tests again on this pod

Clean up

```
kubectl delete deployments.apps webapp-deployment
```

Introduction into istio



@scottcoulton

Introduction into istio

Istio gives us

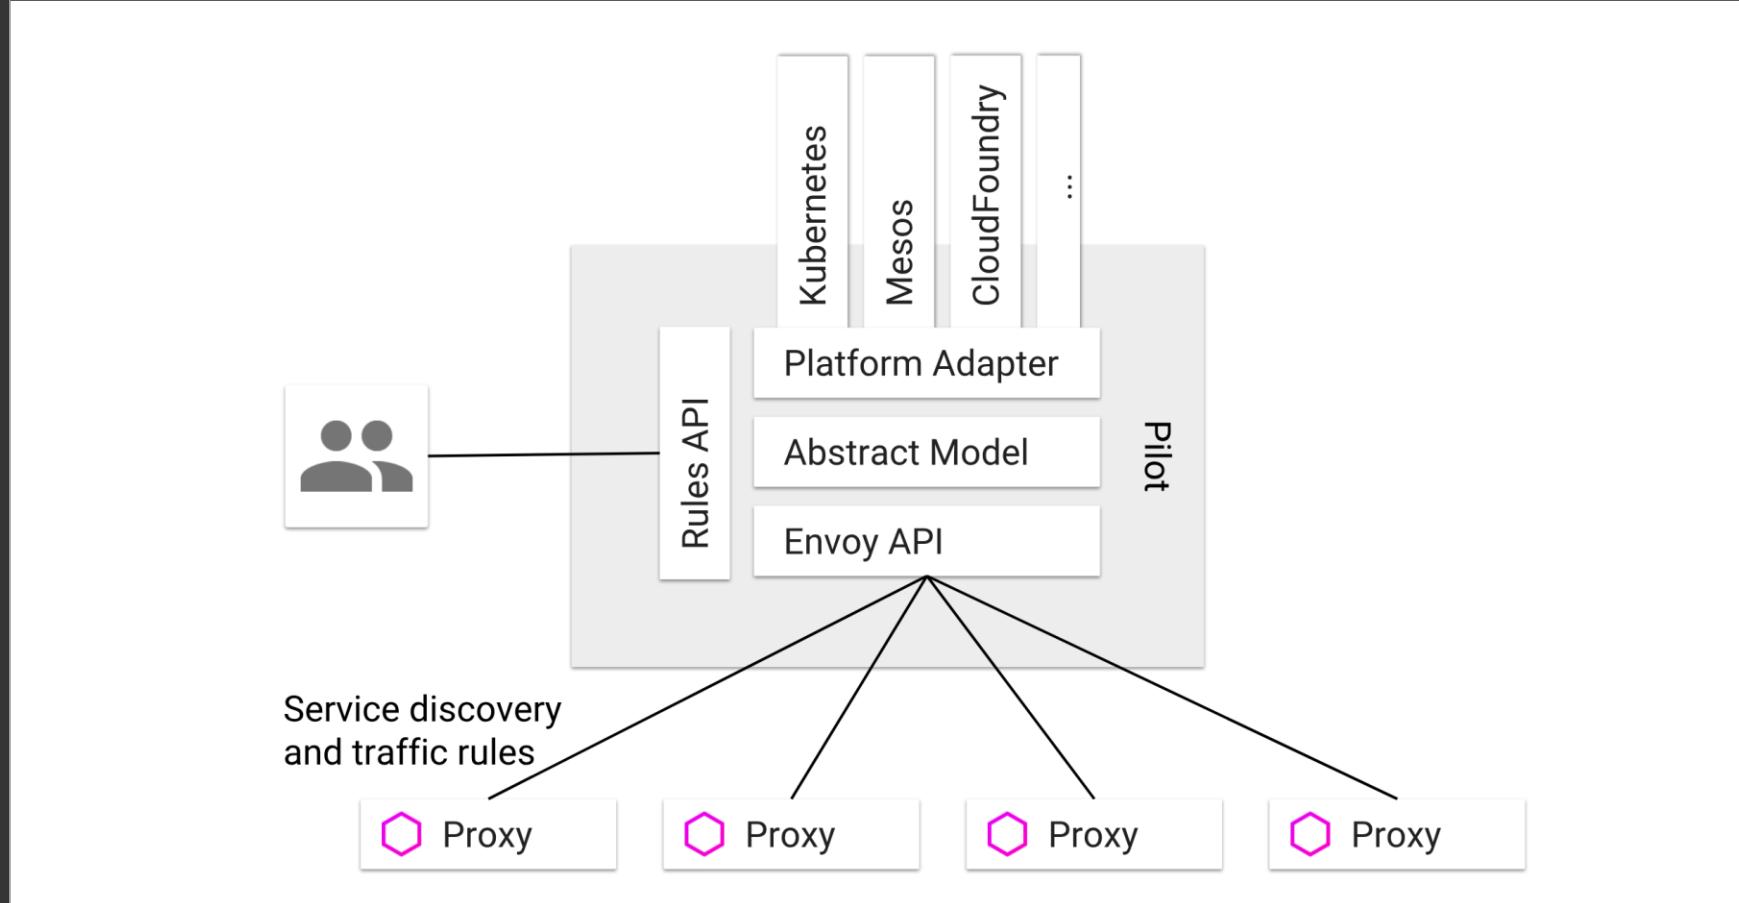
- Automatic load balancing for HTTP, gRPC, WebSocket, and TCP traffic.
- Fine-grained control of traffic behavior with rich routing rules, retries, failovers, and fault injection.
- A pluggable policy layer and configuration API supporting access controls, rate limits and quotas.
- Automatic metrics, logs, and traces for all traffic within a cluster, including cluster ingress and egress.
- Secure service-to-service communication in a cluster with strong identity-based authentication and authorization.

Istio components

Istio is made up of the following components

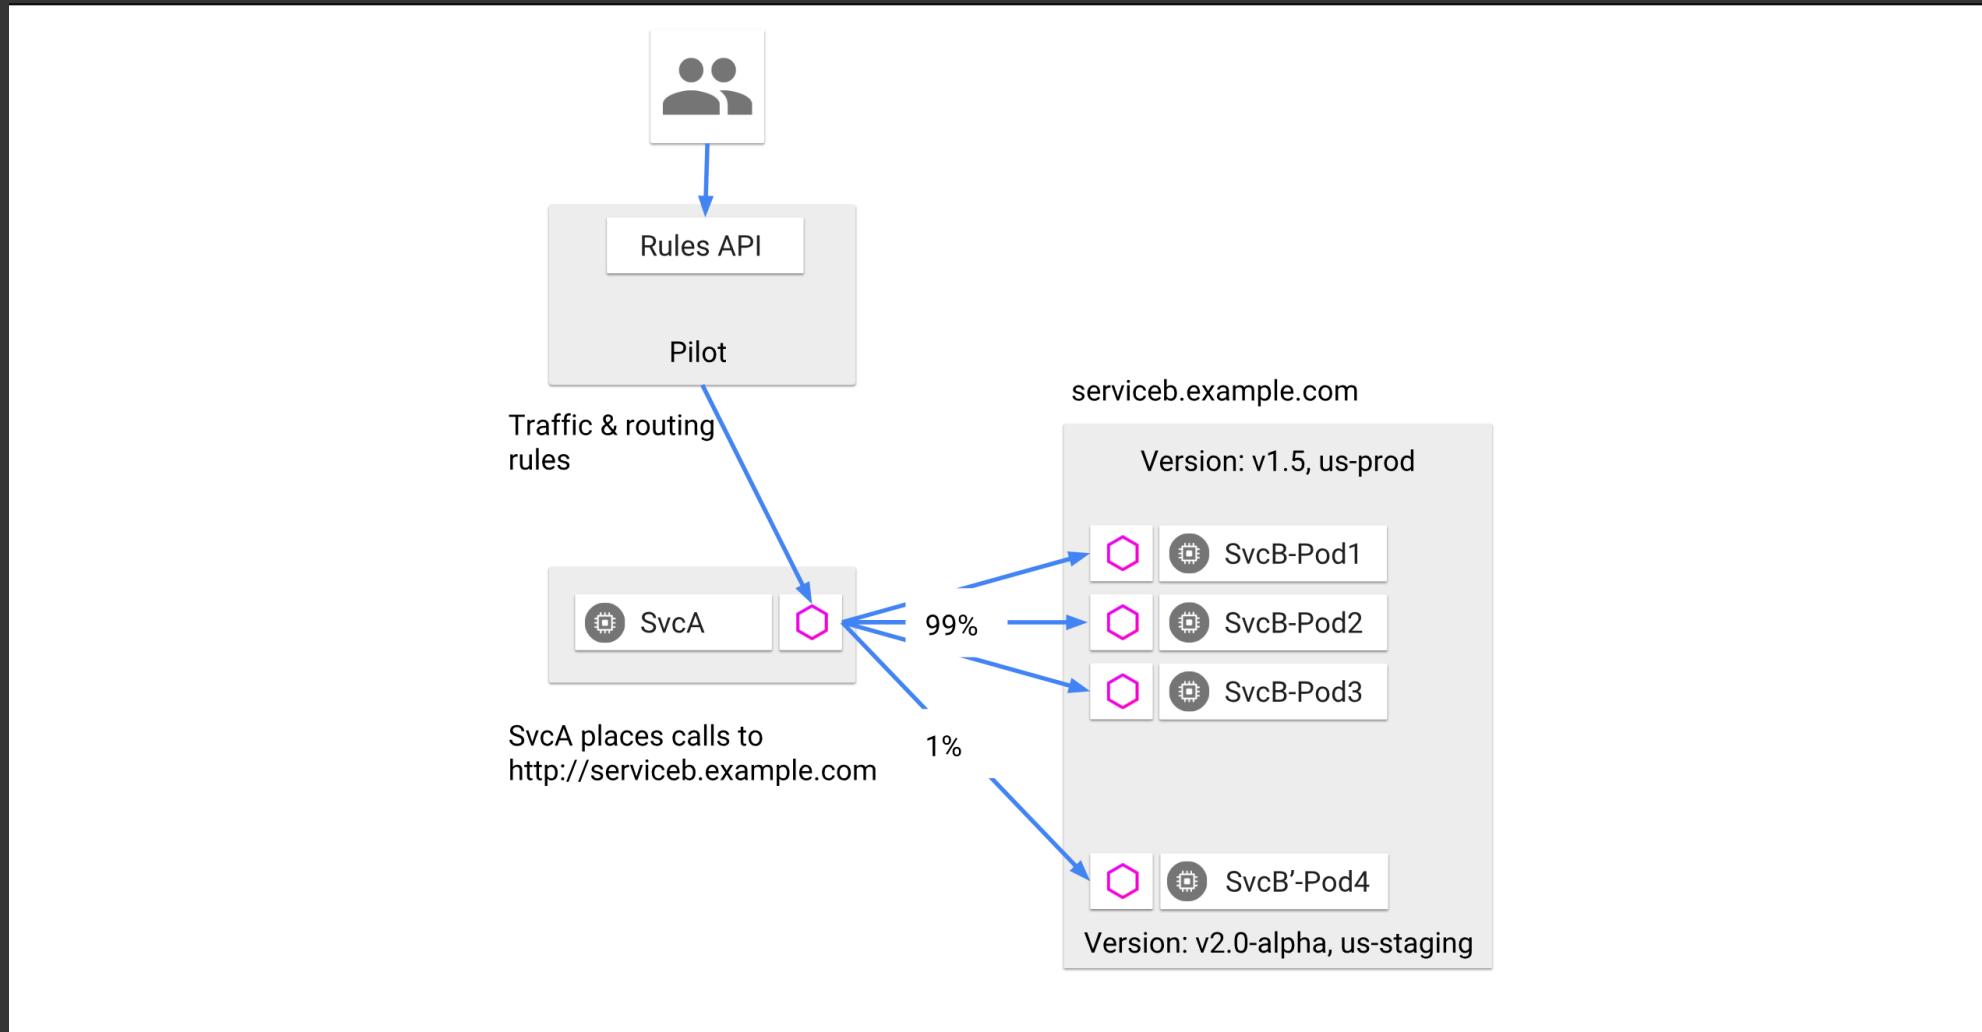
- Envoy and Piolit (ingress, egress & authentication policies)
- Citadel (certificate management)
- Mixer (authorization and auditing)

Envoy and Piolt

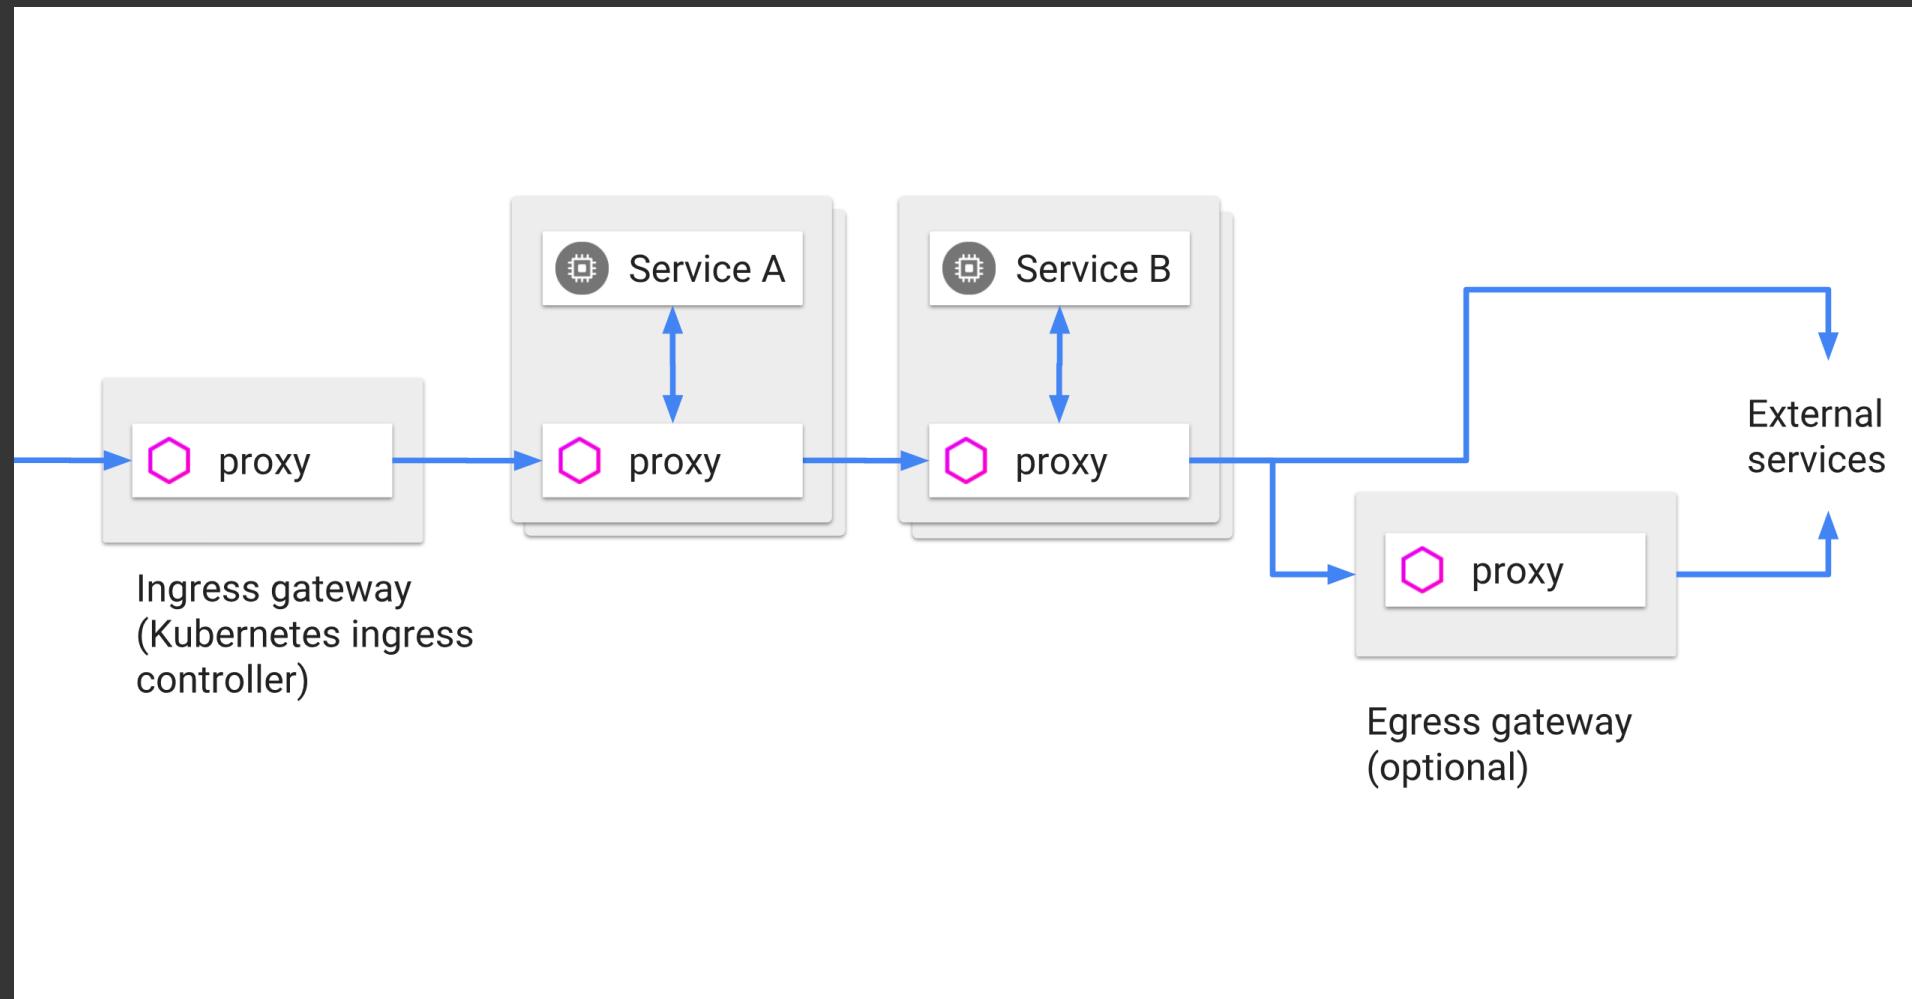


@scottcoulton

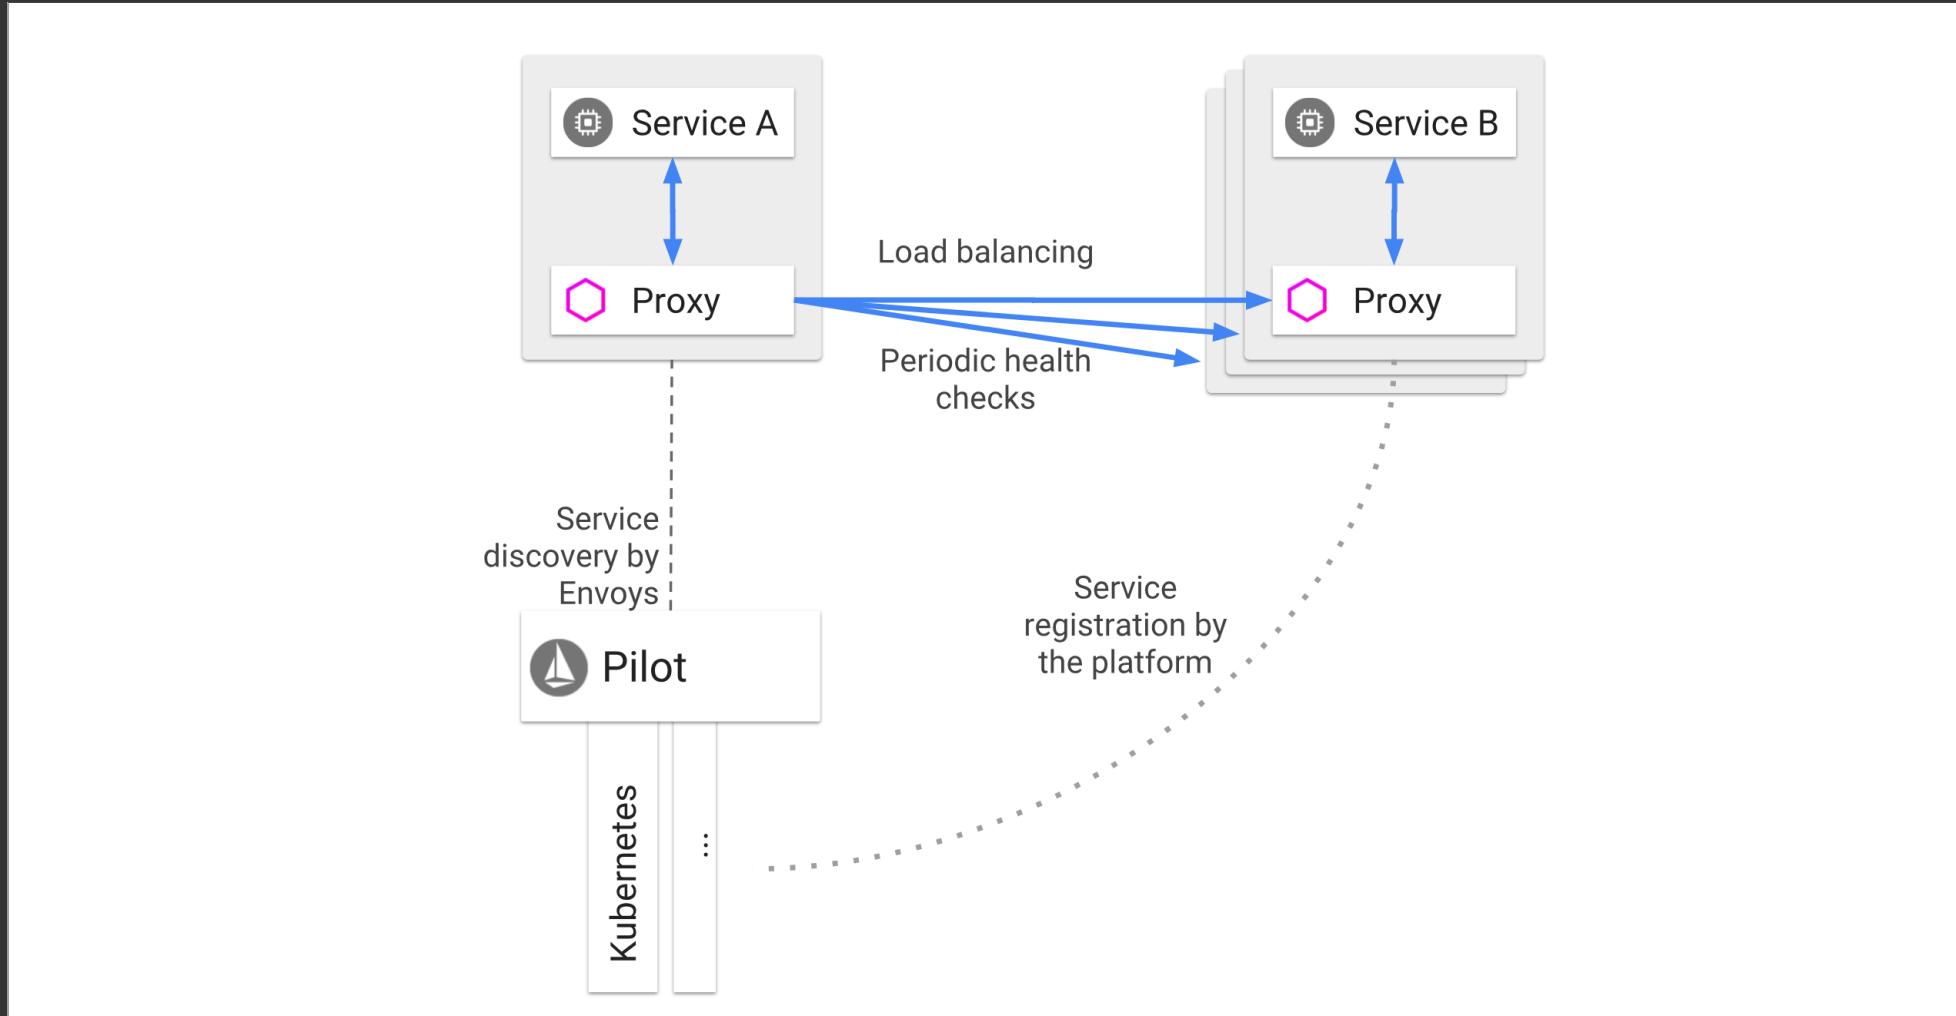
Communication between services



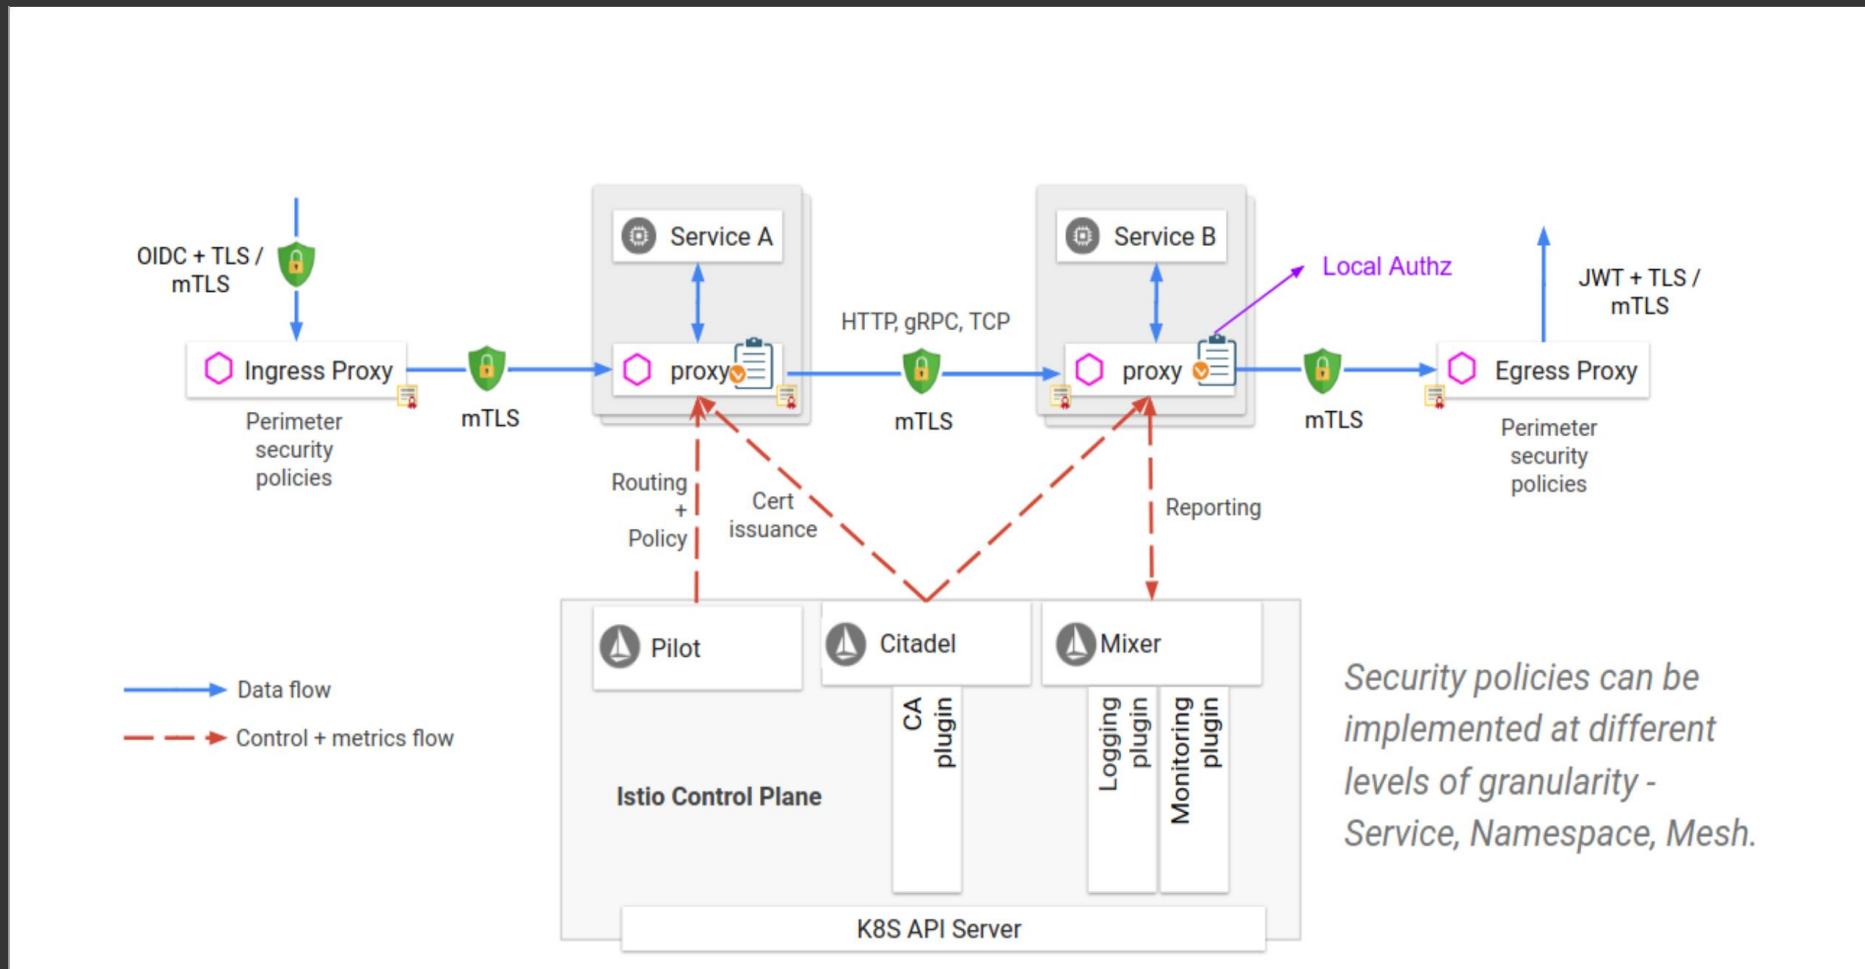
Ingress and egress



Discovery and load balancing



Certificate architecture



@scottcoulton

Advanced application routing with istio



@scottcoulton

Install istio from the script provided

Once istio is installed run

```
kubectl label namespace default istio-injection=enabled
```

Istio application

We will use istio to deploy our webapp

- We will have two versions of the application v1 and v2
- We will weight the traffic evenly between the two
- Now we have the istio crd's installed we will use native kubernetes yaml files to deploy.

Istio application architecture

Add diagram here !!!!

Deploy the services that istio will talk to

```
apiVersion: v1
kind: Service
metadata:
  name: webapp
  labels:
    app: webapp
spec:
  ports:
  - port: 3000
    name: http
  selector:
    app: webapp
```

Create two deployments of the web app

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: webapp-v1
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
        version: v1
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:v1
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 3000
---
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: webapp-v2
spec:
  replicas: 1
  template:
    metadata:
      labels:
        app: webapp
        version: v2
    spec:
      containers:
        - name: webapp
          image: scottyc/webapp:v2
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 3000
```

Create our destination rules

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: webapp
spec:
  host: webapp
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

Create our gateway

```
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: webapp-gateway
spec:
  selector:
    istio: ingressgateway # use istio default controller
  servers:
  - port:
      number: 80
      name: http
      protocol: HTTP
    hosts:
    - "*"
```

Create our virtualservice

```
apiVersion: networking.istio.io/v1alpha3
kind: virtualService
metadata:
  name: webapp
spec:
  hosts:
    - "*"
  gateways:
    - webapp-gateway
  http:
    - route:
        - destination:
            host: webapp
            subset: v1
            weight: 50
        - destination:
            host: webapp
            subset: v2
            weight: 50
```

Get your ingress ip

```
kubectl get svc istio-ingressgateway -n istio-system -o  
jsonpath=".status.loadBalancer.ingress[0].ip"
```

mTLS with istio

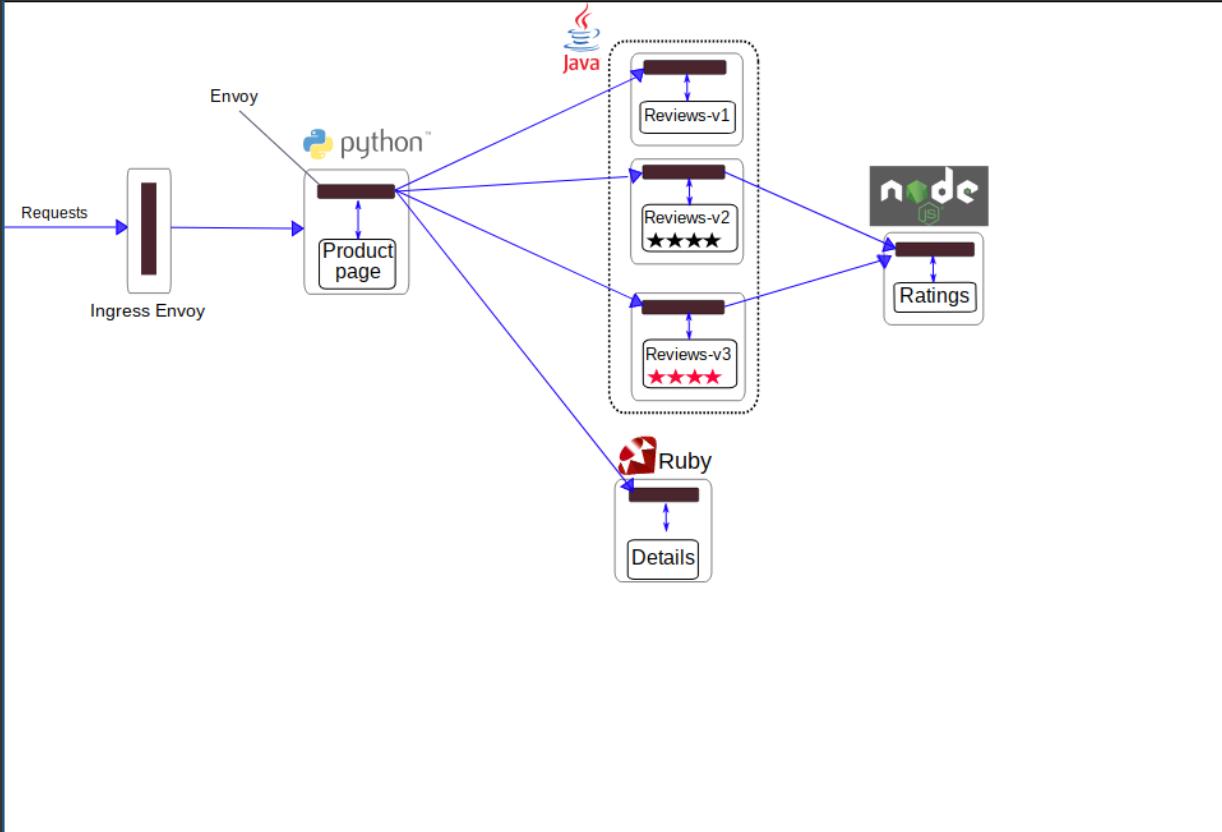


@scottcoulton

Warning !!!!!

mTLS needs to be set up in an empty namespace.

The application we are deploying



@scottcoulton

Create our namespace

```
kubectl create namespace istio-app
```

Make sure we turn on istio injection to our new namespace

```
kubectl label namespace istio-app istio-injection=enabled
```

Enable mTLS across our namespace

```
cat <<EOF | istioctl create -f -
apiVersion: authentication.istio.io/v1alpha1
kind: Policy
metadata:
  name: default
  namespace: istio-app
spec:
  peers:
  - mtls:
EOF
```

Create a destination rule

```
cat <<EOF | istioctl create -f -
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: default
  namespace: istio-app
spec:
  host: "*"
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL
EOF
```

Deploy our application

```
kubectl create -n istio-app -f istio-  
1.0.4/samples/bookinfo/platform/kube/bookinfo  
.yaml
```

Create our virtual service

```
cat <<EOF | kubectl apply -f -
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: bookinfo
  namespace: istio-app
spec:
  gateways:
  - bookinfo-gateway
  hosts:
  - '*'
  http:
  - match:
    - uri:
        exact: /productpage
    - uri:
        exact: /login
    - uri:
        exact: /logout
    - uri:
        prefix: /api/v1/products
  route:
  - destination:
      host: productpage
      port:
        number: 9080
EOF
```

@scottcoulton

Create our gateway

```
cat <<EOF | kubectl apply -f -
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: bookinfo-gateway
  namespace: istio-app
spec:
  selector:
    istio: ingressgateway # use istio default controller
  servers:
  - port:
      number: 80
      name: http
      protocol: HTTP
    hosts:
    - "*"
EOF
```

Create our gateway

```
kubectl -n istio-system get service istio-ingressgateway -o  
jsonpath='{.status.loadBalancer.ingress[0].ip}'
```

We will then use the ip address in our browser to make sure our site is working correctly.

http://<YOUR_IP>/productpage

Check our mTLS status

```
istioctl authn tls-check | grep .istio-  
app.svc.cluster.local
```

@scottcoulton

Check our mTLS status

Exec into our envoy proxy sidecar

```
export POD_NAME=$(kubectl get pods --  
namespace=istio-app | grep details | cut -d' '\"  
-f1)
```

```
kubectl exec -n istio-app -it $POD_NAME -c  
istio-proxy /bin/bash
```

Hit a service internally

```
curl -k -v http://details:9080/details/0
```

@scottcoulton

Let's capture the traffic with tcpdump

```
IP=$(ip addr show eth0 | grep "inet\b" | awk  
'{print $2}' | cut -d/ -f1)  
sudo tcpdump -vvv -A -i eth0 '((dst port 9080)  
and (net $IP))'
```

From another terminal

```
curl -o /dev/null -s -w "%{http_code}\n"  
http://$(kubectl -n istio-system get service  
istio-ingressgateway -o  
jsonpath='{.status.loadBalancer.ingress[0].ip}'  
)productpage
```

Output

```
^[[22:47:36.978639 IP (tos 0x0, ttl 64, id 19003, offset 0, flags [DF], proto TCP (6), length 60)
  10.244.0.12.50662 > details-v1-6764bbc7f7-x7x99.9080: Flags [S], cksum 0x162b (incorrect -> 0xb758), seq 2995501799, win 29200, options [mss 1460,sackOK,TS val 1887650117 ecr 0,nop,wscale 7], length 0
E..<J;@.@...
...
... .#x.....r...+.....
p.AE.....
22:47:36.978681 IP (tos 0x0, ttl 64, id 19004, offset 0, flags [DF], proto TCP (6), length 52)
  10.244.0.12.50662 > details-v1-6764bbc7f7-x7x99.9080: Flags [.], cksum 0x1623 (incorrect -> 0x3e8d), seq 2995501800, ack 2809488904, win 229, options [nop,nop,TS val 1887650117 ecr 2183432464], length 0
E..4J<@.@...
...
... .#x....uf.....#....
p.AE.$..
22:47:36.978742 IP (tos 0x0, ttl 64, id 19005, offset 0, flags [DF], proto TCP (6), length 254)
  10.244.0.12.50662 > details-v1-6764bbc7f7-x7x99.9080: Flags [P.], cksum 0x16ed (incorrect -> 0xe838), seq 0:202, ack 1, win 229, options [nop,nop,TS val 1887650117 ecr 2183432464], length 202
E...J=@.@...
...
... .#x....uf.....
p.AE.$.....A.....}.Q..d.....+.../. ....../.,.0.
....5...|.....7.5..2outbound|9080||details.istio-app.svc.cluster.local.....#.....istio.....
.....
22:47:36.979801 IP (tos 0x0, ttl 64, id 19006, offset 0, flags [DF], proto TCP (6), length 52)
  10.244.0.12.50662 > details-v1-6764bbc7f7-x7x99.9080: Flags [.], cksum 0x1623 (incorrect -> 0x38ac), seq 202, ack 1280, win 251, options [nop,nop,TS val 1887650118 ecr 2183432465], length 0
E..4J>@.@...
...
... .#x....uk.....#....
p.AF.$..
22:47:36.981099 IP (tos 0x0, ttl 64, id 19007, offset 0, flags [DF], proto TCP (6), length 1245)
  10.244.0.12.50662 > details-v1-6764bbc7f7-x7x99.9080: Flags [P.], cksum 0x1acc (incorrect -> 0xba00), seq 202:1395, ack 1280, win 251, options [nop,nop,TS val 1887650120 ecr 2183432465], length 1193
E...J?@.@...
...
... .#x....uk.....
p.AH.$.....6..3..00..,0.....->.....).....0... *.H.....0.1.0...U.
..k8s.cluster.local0...190108215913Z..190408215913Z0.1 0...U.
...
... .#x....up....7.#....
p.AM.$..
```



Azure ❤️ 🛡️

